

Operating Systems Lab: Assignment 1

Members: Kousik Rajesh, Drishti Chouhan, Eklavya Jain, Mridul Garg

Q.1 . asm("inc %0": "+r"(x));

```
kousik@sloth:~/xv6-public$ gcc ex1.c
kousik@sloth:~/xv6-public$ ./a.out
Hello x = 1
Hello x = 2 after increment
OK
```

The **inc** stands for incrementing the value of the provided argument. **"+r"** stands for the fact that **x** is the input as well as the output. And **%0** denotes the first and the only argument, **x**.

Q.2

```
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: ljmp $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
+ symbol-file kernel
(gdb) si
[f000:e05b] 0xfe05b: cmpw $0xffc8,%cs:(%esi)
0x0000e05b in ?? ()
(gdb) si 2
[f000:e066] 0xfe066: xor %edx,%edx
0x0000e066 in ?? ()
(gdb) si
[f000:e068] 0xfe068: mov %edx,%ss
0x0000e068 in ?? ()
(gdb) si
[f000:e06a] 0xfe06a: mov $0x7000,%sp
0x0000e06a in ?? ()
(gdb) si
[f000:e070] 0xfe070: mov $0x7c4,%dx
0x0000e070 in ?? ()
(gdb) si
[f000:e076] 0xfe076: jmp 0x5576cf26
0x0000e076 in ?? ()
(gdb) si
[f000:cf24] 0xfc24: cli
0x0000cf24 in ?? ()
(gdb) si
[f000:cf25] 0xfc25: cld
0x0000cf25 in ?? ()
(gdb) si
[f000:cf26] 0xfc26: mov %ax,%cx
0x0000cf26 in ?? ()
(gdb) si
[f000:cf29] 0xfc29: mov $0x8f,%ax
0x0000cf29 in ?? ()
(gdb) si
[f000:cf2f] 0xfc2f: out %al,$0x70
0x0000cf2f in ?? ()
(gdb) si
[f000:cf31] 0xfc31: in $0x71,%al
0x0000cf31 in ?? ()
(gdb) si
[f000:cf33] 0xfc33: in $0x92,%al
0x0000cf33 in ?? ()
```

```
[f000:cf35] 0xfc35: or $0x2,%al
0x0000cf35 in ?? ()
(gdb) si
[f000:cf37] 0xfc37: out %al,$0x92
0x0000cf37 in ?? ()
(gdb) si
[f000:cf39] 0xfc39: mov %cx,%ax
0x0000cf39 in ?? ()
(gdb) si
[f000:cf3c] 0xfc3c: lidt %cs:(%esi)
0x0000cf3c in ?? ()
(gdb) si
[f000:cf42] 0xfc42: lgdtl %cs:(%esi)
0x0000cf42 in ?? ()
(gdb) si
[f000:cf48] 0xfc48: mov %cr0,%ecx
0x0000cf48 in ?? ()
(gdb) si
[f000:cf4b] 0xfc4b: and $0xffff,%cx
0x0000cf4b in ?? ()
(gdb) si
[f000:cf52] 0xfc52: or $0x1,%cx
0x0000cf52 in ?? ()
(gdb) si
[f000:cf56] 0xfc56: mov %ecx,%cr0
0x0000cf56 in ?? ()
(gdb) si
[f000:cf59] 0xfc59: ljmpw $0xf,$0xfc61
0x0000cf59 in ?? ()
(gdb) si
The target architecture is assumed to be i386
=> 0xfc61: mov $0x10,%ecx
0x0000fc61 in ?? ()
(gdb) si
=> 0xfc66: mov %ecx,%ds
0x0000fc66 in ?? ()
(gdb) si
=> 0xfc68: mov %ecx,%es
0x0000fc68 in ?? ()
(gdb) si
=> 0xfc6a: mov %ecx,%ss
0x0000fc6a in ?? ()
```

- The BIOS starts with the assumption that the architecture is i8086(1MB address space). We first switch the mode to i386 by making a jump to a previous address
[f000:fff0] 0xffff0: ljmp \$0x3630:0xf000e05b
- The BIOS then zeros the **edx** and **ss** registers and sets **sp** register and configures the hardware devices using the **out** and **in** command
- The **cli instruction** disables interrupts so that the following instructions will be executed without interruption. It does so by clearing the interrupt flag
- The **cld instruction** clears the direction flag and sets it to zero indicating that memory grows from low to high
- Finally, it loads the boot sector at the address **0x7c00**

Q3. a)

```
# Switch from real to protected mode
lgdt     gdt_desc
movl     %cr0, %eax
orl      $CR0_PE, %eax
movl     %eax, %cr0
# Following long jmp completes the transition to 32-bit protected mode
ljmp     $(SEG_KCODE<<3), $start32#
```

b) Last instruction that Bootloader executes :

In **bootmain.c** it is where the **entry** function is called to enter the kernel

```
entry = (void (*)(void))(elf->entry);
entry();
```

In **bootblock.asm** it is **call *0x10018** which calls the entry function

```
7d87: ff 15 18 00 01 00 call *0x10018
```

First instruction of the kernel :

```
movl     %cr4, %eax
```

(The first instruction is present at **0x0010000c**)

```
(gdb) b * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/1x 0x10018
0x10018: 0x00000000
(gdb) b * 0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d91: call *0x10018

Thread 1 hit Breakpoint 2, 0x00007d91 in ?? ()
(gdb) si
=> 0x10000c: mov     %cr4,%eax
0x0010000c in ?? ()
(gdb) x/1x 0x10018
0x10018: 0x0010000c
```

c) The code segment which loads the kernel sectors from disk

```
ph = (struct prog_hdr *) ((uchar *)elf + elf->phoff);
eph = ph + elf->phnum;
for (; ph < eph; ph++) {
    pa = (uchar *)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
}
```

- **ph** Points to the start of the program header table found by adding an offset to **elf**
- **elf->phnum** contains the number of entries in the program header table.
- **eph** is a pointer to the entry just after the last entry in program header table
- We iterate through each sector and read them one by one by incrementing the pointer **ph** till we reach **eph** at which point we exit the for loop

Q.5 On changing the link address in Makefile from **0x7C00** to **0x7D00** the code expected to be at **0x7C00** is not present

The first instruction to go wrong is

```
ljmp     $0x3630, $0xf000e05b
```

The bootloader is stuck at booting from hard disk. When inspected using **GDB** I found that the **IP(instruction pointer)** begins to loop between address **6dc1** and **eeee** it overflows at **eeee** and then goes back to **6dc1** in an infinite loop.

```
SeaBIOS (version 1.13.0-1ubuntu1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

Booting from Hard Disk..
```

Q.6

When BIOS enters Boot loader

All are zero

Doing an **objdump** for kernel we can see that it's

Load Memory Address(LMA) is at 0x00100000

And since the kernel hasn't been loaded yet, these 8 words in memory are zero.

When Boot loader enters Kernel

The bootloader has now finished reading all sectors from disk using the **readseg()** function and the address **0x00100000** falls in the **.text** section of the kernel and hence when we now examine the memory near this address we find the kernel instructions which were newly loaded by the bootloader

```
kousik@sl0th:~/xv6-public$ objdump -h kernel

kernel:      file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
  0 .text          0000713a  80100000  00100000  00001000  2**4
CONTENTS, ALLOC, LOAD, READONLY, CODE

(gdb) b * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) b * 0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x7d91: call *0x10018

Thread 1 hit Breakpoint 2, 0x00007d91 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
(gdb) x/8i 0x00100000
0x100000: add    0x1bad(%eax),%dh
0x100006: add    %al,(%eax)
0x100008: decb   0x52(%edi)
0x10000b: in     $0xf,%al
0x10000d: and    %ah,%al
0x10000f: or     $0x10,%eax
0x100012: mov    %eax,%cr4
0x100015: mov    $0x109000,%eax
(gdb) █
```

Q.7 Files changed:

sysfile.c : Add the system call which copies wolfie into the provided buffer and returns an integer

user.h : Add a function declaration for the system call in the correct format

usys.S : Add a SYSCALL() for wolfie

syscall.h : Define a number for the wolfie syscall

syscall.c : Add our custom defined syscall to the list of syscalls present in this file

Q.8 Files changed:

wolfietest.c : This will be our user level application which uses the syscall and prints wolfie on the terminal

Makefile : Inform the Makefile compile our user level application wolfietest.c

\$ **■**